

ENTERPRISE AI PRODUCTION · AGENT SKILLS & PRODUCT LIBRARY

Patrick Diamitani

GTM AI & Automation Manager building production AI systems that 50+ non-technical revenue teammates use every day — including a customer-facing AI agent that has already sourced a quarter-million-dollar deal and \$520K+ in attributed pipeline.

**14**

Systems in the library

9

Live in production

50+

Daily non-technical users

[View the library ↓](#)[About Patrick](#)

About

Patrick Diamitani owns AI strategy and automation across an entire go-to-market organization — sales, marketing, and revenue operations — at a global HR-tech (Employer of Record) platform with roughly 500 employees operating in 160+ countries. His mandate: build production systems that a non-technical team of 50+ reps and marketers can rely on daily, not demos that impress once and then quietly stop being used.

Four themes run through everything in this library. **Hybrid AI architecture** — large language models handle what they're good at (reading messy text, drafting, classification), while deterministic, unit-tested code handles anything that must be auditable, like compliance tiers or money. The system never just "lets the AI decide" on a consequential call. **Production discipline** — every build ships with dry-run tests, rollback plans, and shared error handling, because these systems run unattended on live revenue pipelines, not in a sandbox. **Cost engineering** — AI and data-enrichment spend gets measured and capped by design, not by asking people to be careful. **Platform building** — beyond any single automation, the real product is a reusable skill library, a project-intake system, and a build standard so the work compounds instead of resetting with every new request.

The result is a portfolio that reads less like a collection of scripts and more like a small internal product organization: an AI agent with its own attribution model, a compliance engine with full audit trails, a cost-governance framework that turned an owned mistake into a permanent safeguard, and the standards layer that let one person's work scale across fifty.



Career Highlights

- Built and shipped a full-stack AI customer agent — chat interface, qualification pipeline, reporting warehouse, and revenue-attribution layer — that sourced a \$250K deal traceable to a single conversation and \$520K+ in attributed pipeline within its first four months live.
- Designed and launched an automated compliance-screening system that corrected 100+ misclassified deals stuck in manual review and more than doubled detection of ineligible deals across the full sales pipeline, with a perfect production write record.
- Diagnosed and rebuilt a four-agent, post-sales-call automation pipeline from a failing RAG-based design (a broken vector index causing the majority of failures) to a deterministic architecture — going from 0 successful end-to-end runs to fully verified in production.
- Owned a runaway AI/data-enrichment spend incident end to end and redesigned the underlying system around a free eligibility gate and tiered cost profiles, cutting spend per record by roughly 70% by architecture, not vigilance.
- Directed a nine-script enrichment pipeline for an outsourced sales-development pipeline, achieving confirmed 0% contact overlap against thousands of already-contacted records and 91% phone coverage.
- Built and published a 40+ skill AI automation library plus a self-serve project-intake engine that cut idea-to-scoped-project time from 4–8 hours to under 30 minutes across a 50-person go-to-market org.

Approach

Hybrid AI Architecture

Every system in this library draws a hard line between judgment and interpretation. Language models are used for exactly what they're good at — reading unstructured text, mapping messy values to clean categories, drafting first-pass copy — while the actual decision (a compliance tier, a lifecycle stamp, a dollar figure) is made by deterministic, testable code that produces the same answer on the same input, every time.

Production Discipline

Nothing ships without a dry-run path, a rollback plan, and a way to fail loudly instead of silently. Systems that touch live CRM data default to report-only until a human signs off, and every automation shares a common error-handling pattern so a failure surfaces the same way, every time, instead of becoming a one-off mystery.



Cost Engineering

AI and enrichment spend is treated as an engineering constraint, not an afterthought. Free eligibility checks run before any paid step, expensive operations are gated to the accounts that actually warrant them, and every workflow reports its own cost — so spend stays predictable by design rather than by someone remembering to watch a dashboard.

Platform Building

The most durable output isn't any single automation — it's the standard that makes the next one faster to build safely. A shared build methodology, a self-serve project-intake tool, and a reusable skill library mean the organization's AI capability compounds instead of living entirely inside one person's head.

Skill & Product Library

Search the library...

- All (14)
- Compliance
- Data Quality
- Observability
- Platform
- RevOps
- Sales Automation

PlatformLive in production

The GTM AI Platform

One automation is a win; fifty people using AI daily requires a system. A published build methodology, a self-serve project-intake engine, and a 40+ skill library now let one person's work compound across an entire go-to-market org instead of resetting with every new request.

Standards, governance, and scaleRead the build →

AI Agent SystemLive in production

AI Customer Agent Suite

A global HR-tech (EOR) platform needed an AI agent that could qualify website visitors and prove its own revenue impact, not just chat. Built end to end — agent, qualification pipeline, reporting warehouse, and a three-tier attribution model — it sourced a \$250K deal traceable to a single conversation and \$520K+ in attributed pipeline within four months.

Full-stack AI product ownershipRead the build →

AI Agent SystemLive in production

Contract Intake Automation (DocuSign)

Sellers were hand-copying deal data into a legal contract-intake form, and any typo stalled the deal. After two 'obvious' API integrations were proven blocked by the vendor itself, the automation pivoted to the one door left open — bringing intake time from a manual chore down to under 3 minutes.

Constraint-driven architectureRead the build →

Automation WorkflowLive in production

EOR Compliance Checker

Deals were reaching contracting before anyone checked whether the role or industry was even eligible for the company's core service, killing deals late. An LLM-plus-rules-engine screening system now evaluates every deal on entry, correctly clearing 100+ deals stuck in manual review and doubling detection of ineligible deals with a perfect production write record.

Regulated-domain AI with audit trailsRead the build →



Claude Skill

Live in production

HubSpot Import Formatter

Inconsistent rep-uploaded spreadsheets were polluting the CRM with duplicate contacts and companies, inflating pipeline forecasts. A skill now converts any spreadsheet into clean, deduped, NAICS-enriched import templates, live across a 50+ person team with adoption tracked by a run ledger.

Root-cause data-quality engineering

[Read the build →](#)

Automation Workflow

Live in production

Partner PQL Automation

A partner-inquiry form mixed genuine partnership leads with vendor spam, and everything landed on one person to triage by hand. An LLM now classifies every submission and routes only real leads to a high-priority task with an AI-drafted response, filtering spam with zero manual effort.

LLM classification in production CRM

[Read the build →](#)

Automation Workflow

Live in production

Post-Sales Call Automation

Reps left sales calls without consistent follow-up, and country-specific facts — the core of this company's product — varied rep to rep. A four-agent pipeline now turns every recorded call into five deliverables automatically, rebuilt from a failing RAG-based design (0/25 passing) to a fully verified deterministic one.

Event-driven multi-agent pipelines [Read the build →](#)

Data Pipeline

Live in production

Prospect Enrichment Pipeline

An outsourced sales-development team was being fed stale, already-called contacts, burning paid outreach time. A nine-script Python pipeline now merges, dedupes, and enriches prospect lists automatically, delivering confirmed 0% overlap against thousands of already-contacted records and 91% phone coverage.

API-based enrichment at scale [Read the build →](#)

AI Agent System

Phase 1 live

RFP Response Suite

Enterprise RFP responses took days of document archaeology, and knowledge gaps stayed invisible until a prospect's question exposed them. A three-part skill suite now builds an atomic, owner-routed knowledge base and auto-creates a tracked ticket for every

Platform

Shipped

Clay Credit Governance

A runaway enrichment-spend incident burned through a large credit purchase in about a month with zero pre-run cost visibility. Owning the failure directly led to a rebuilt, tiered enrichment standard with a free eligibility gate and per-row cost estimates — cutting spend per record by roughly 70% by design.

unanswered question instead of letting it disappear.

Knowledge-base architecture

[Read the build →](#)

FinOps for AI/data tooling

[Read the build →](#)

Claude Skill

Shipped

Clay Prospecting Engine

A powerful enrichment platform's list-building UI was never something 50+ non-technical reps could safely drive themselves, so every request bottlenecked on one person. A CSV-upload skill now auto-detects company vs. contact lists and routes to the right enrichment pipeline with zero UI interaction required.

Productizing internal tooling

[Read the build →](#)

Automation Workflow

Build-ready

Entity Readiness Change Report

A company's legal ability to hire in a given country changes over time, but nobody was catching the flips — marketing kept paying for ads in countries that were no longer eligible. A failure-safe weekly diff engine now reports only what changed, designed so a bad data pull can never corrupt the system's memory.

State management & idempotency

[Read the build →](#)



Automation Workflow

Shipped

Pre-Call Briefing Automation

Reps walked into discovery calls cold despite the context already existing somewhere in the CRM. This workflow assembles contact, deal, and buyer-intent data automatically and delivers an AI-drafted brief to the rep's inbox before every meeting, with built-in dedupe so nothing is briefed twice.

Sales-call lifecycle ownership

[Read the build →](#)

Platform

Design phase

n8n Execution Intelligence

When an automation failed or quietly stalled, the team found out only when data stopped moving, then diagnosed it by clicking through runs one at a time. A read-only observability layer now targets cutting diagnosis time from 10-15 minutes to under 60 seconds, with a hard no-fabrication charter on every claim.

SRE instincts for automation

[Read the build →](#)

Let's talk

Reach Patrick at pdiamitani@atlashxm.com or on LinkedIn at linkedin.com/in/diamitani.

[Email](mailto:pdiamitani@atlashxm.com)[LinkedIn ↗](https://linkedin.com/in/diamitani)

© 2026 Patrick Diamitani. Case studies anonymized for public portfolio use — no employer-identifying or client-identifying detail included.

pdiamitani@atlashxm.com · linkedin.com/in/diamitani