

Skill & Product Library / RevOps

AI Agent System

Phase 1 live

RFP Response Suite

From days to hours on enterprise questionnaires — a knowledge base that closes its own gaps.

Phase 1 live: atomic KB + auto-created HubSpot tickets for every answer gap



Claude Skill

HubSpot

RFP

knowledge base

semantic search

Overview

Enterprise RFP responses at a global HR-tech (EOR) platform took days of hunting through old documents, answers were inconsistent between reps, and gaps in the knowledge base were invisible until a prospect asked a question nobody could answer. This three-part Claude skill suite fixes all three: an **RFP KB Builder** parses historical RFP files and support tickets into atomic, owner-routed Q&A facts; an **RFP Answer Retriever** semantically matches new RFP questions against that knowledge base with an honest confidence rating and auto-creates tickets for every gap; and an **RFP Responder agent** ingests a full RFP, categorizes every requirement, and drafts compliant responses. Phase 1 (KB + gap tickets) is live; Phases 2–3 (compliance-platform integration, auto-draft) are spec'ed and ready to build.

The Problem

RFP responses took days of manual document archaeology, answers drifted between reps answering the same question differently over time, and there was no visibility into which facts the company simply didn't have a good answer for — gaps surfaced only when a prospect's question stumped someone.

How It Works

-

Tool	Role	Access needed
Claude skills + agent	KB building, retrieval, drafting	Claude Code / Claude with skills installed
Document parsing (PDF/DOCX/XLSX)	Ingests historical RFPs and source documents	Local file access
Semantic matching engine	Confidence-tiered retrieval against the KB	Ships with the Answer Retriever skill
CRM tickets API (HubSpot)	Auto-creates a ticket for every knowledge gap	Private-app token (tickets read-write)

Tool	Role	Access needed
Excel export	Two-tab, compliance-platform-ready KB workbook	Local file output
GRC platform (Onspring, later phase)	Target destination for a live-synced KB	Planned integration

Setup Process

1. Gather historical RFP files and export relevant support tickets as source material.
2. Run the RFP KB Builder skill to generate the atomic Q&A knowledge base; review and resolve any flagged conflicts.
3. Assign real owners (product/HR services/finance/global mobility/implementation) to the fact categories in your org.
4. Export the two-tab Excel workbook; this is your live KB for Phase 1.
5. Run the RFP Answer Retriever skill against a new RFP; review the coverage summary and confirm gap tickets land correctly in the CRM.
6. (Later phases, when ready) Wire the KB to your GRC platform and enable the full RFP Responder agent for auto-drafted responses.



What You Need

- [] Historical RFP documents (.pdf/.docx/.xlsx) and relevant CRM ticket history
- [] Claude Code or Claude access with the RFP skill suite installed
- [] CRM private-app token (tickets read-write) for auto-gap-ticketing
- [] Named fact owners across product/HR services/finance/global mobility/implementation
- [] ~1–2 days for initial KB build; hours per RFP thereafter

Results

- **Phase 1 live:** Excel-as-active-KB plus automatic CRM ticket creation for every answer gap.
- **Later phases speccd:** GRC platform integration, confidence-scoring v2, and full auto-draft generation.
- Every gap becomes an owned, tracked ticket instead of disappearing until the next RFP hits the same wall.

Challenges & Fixes

- **Answers drifting between reps over time** → atomic Q&A facts with named owners instead of free-text institutional memory.
- **Region-specific facts misapplied company-wide (or vice versa)** → explicit general vs. region-specific classification on every fact.
- **Historical answers contradicting each other** → conflicts flagged for human resolution instead of silently auto-resolved.
- **Knowledge gaps invisible until a prospect asked** → NO-MATCH questions auto-create a CRM ticket, closing the loop from "unknown" to "assigned."

Demonstrates

Knowledge-base architecture (atomic facts, ownership routing, conflict detection) · honest-confidence retrieval design (HIGH/MEDIUM/NO-MATCH, never a bluff) · closing the loop from "unknown" to "assigned."



STACK & INTEGRATIONS

Claude Skill

HubSpot

Excel

Onspring (planned)

WHAT IT NEEDS

- ✓ Historical RFP documents (.pdf/.docx/.xlsx) and CRM ticket history
- ✓ Claude Code or Claude access with the RFP skill suite
- ✓ CRM private-app token (tickets read-write)
- ✓ Named fact owners across product/HR/finance/mobility/implementation

DISCUSS THIS BUILD

Happy to walk through the architecture, the failures along the way, or how this could translate to your team.

Email Patrick

© 2026 Patrick Diamitani. Case studies anonymized for public portfolio use — no employer-identifying or client-identifying detail included.

pdiamitani@atlashxm.com · [linkedin.com/in/diamitani](https://www.linkedin.com/in/diamitani)

