

Skill & Product Library / Platform

Platform

Live in production

The GTM AI Platform

Not one automation — the operating system a 50-person GTM org runs its AI adoption on.

40+ skills live across a 50-person org; idea-to-scoped-project cut from 4-8 hrs to under 30 min



Claude Skills

n8n

Ollama

platform

governance

ROSTR

Overview

Any single automation is a win. Fifty-plus reps and marketers using AI daily requires a *system* — standards, reusable skills, project governance, and knowledge that survives its author. This is the platform layer every other product in this portfolio runs on: a published **build methodology** (a framework applied to every agent in the org), a **project-intake engine** that turns any raw idea into a leadership-ready package in under 30 minutes, a **40+ skill library** covering brand content, GTM automation, chatbot management, and cross-session memory, and a **self-hosted local AI infrastructure** for zero-API-key development. This is the Director-level story: not automations, but the standards, governance, and observability layer that lets one person's work compound across fifty.

The Problem

One-off automations don't scale past their author. Without shared standards, every new skill reinvents its own conventions, every new project idea takes hours of back-and-forth to scope, institutional knowledge lives in one person's head, and there's no way for a non-technical org of 50+ to safely adopt AI at the pace the business needs.

How It Works

1. **A published build methodology.** A framework combining **prompt compilation** (setting technical depth and framing for a given audience), **jobs-to-be-done planning** (what outcome the build actually serves), a **work-classification spine** (navigate/prioritize/allocate/orchestrate — a repeatable shape for any project), and **provenance honesty** (never presenting an inference as a verified fact). Applied to every agent manifest built across the platform, so a new skill isn't a blank page — it inherits a proven shape.

2. **A project-intake engine.** Takes any raw idea — text, files, links, a chat thread — and turns it into a leadership-ready package: a PRD, an architecture diagram, a build plan, a KPI framework, and a project-management export. Cut idea-to-scoped-project from 4–8 hours of manual scoping to **under 30 minutes**, and made every resulting project auditable by anyone, not just its author.

3. **A 40+ skill library for the GTM team.** A brand suite (on-brand social/email/blog/proposal generation using official fonts, colors, and logos — so content stays on-brand without a designer in the loop for every post); a proposal & presentation builder (branded slide-deck output); a GTM automation architect skill (enrichment/workflow/CRM automation expert, callable in plain English); a workflow architect & engineer (turns a rough automation idea into a full technical spec); a chatbot manager; a context-engine (persistent memory across AI sessions, so context doesn't reset every conversation); the case-study builder that generated this very document; a token-spend analyst; and more — each a reusable, packaged capability rather than a one-off script.

4. **Local AI infrastructure.** A self-hosted stack running an AI coding assistant against local open models via a custom translation proxy — a small Python service that lets the assistant's API calls target locally-run models instead of a hosted API. Persisted as background services (surviving reboots without manual restarting) with a model-routing config, enabling zero-API-key development and cost-controlled experimentation before anything touches production credentials.

5. **Governance that compounds.** Every skill inherits the same production discipline visible across this portfolio: dry-run defaults, audit trails, ownership routing, and provenance honesty — so the system gets *safer* as it gets bigger, not riskier.
-
- Stack & Integrations
- | Tool | Role | Access needed |
|--|---|---|
| Claude / Anthropic ecosystem (skills, MCP servers, agents) | The runtime every skill and agent in the platform is built on | Claude Code / Claude access, org-level skill deployment |
| Python | Intake engine logic, local proxy service, analysis pipelines | Local/sandboxed runtime |
| | | |
- <https://diamitani.github.io/patrick-diamitani-portfolio/products/gtm-ai-platform.html>

2/5

Tool	Role	Access needed
n8n	Orchestration layer for automation-heavy skills	n8n workspace
Background service manager (e.g. launchd)	Persists local proxy + model-runtime services across reboots	Local machine admin
Local model runtime (e.g. Ollama)	Local model runtime for zero-API-key development	Local install
GTM SaaS stack (CRM, enrichment, sequencing, call-intelligence, project management, intent data, productivity suite)	The systems every downstream skill in the platform integrates with	Per-tool credentials, scoped per skill



Setup Process

1. Stand up an AI coding/skill environment as the org's shared skill-execution runtime.
2. Publish a build methodology spec as the required shape for any new skill or agent build.
3. Deploy a project-intake skill as the front door for any new automation idea — text, file, or link in, leadership package out.
4. Roll out the skill library in phases (brand suite first for immediate content value, then automation-facing skills).
5. Stand up a context-engine for persistent cross-session memory once the team is running enough recurring work to benefit from it.
6. Optionally stand up local AI infrastructure (local model + translation proxy + background services) for cost-controlled experimentation separate from production API spend.
7. Establish a lightweight review cadence so the skill library doesn't drift from the standard.

What You Need

- [] AI coding/skill environment access for the org (or at minimum, for the platform owner)
- [] A written intake/PRD convention the org will actually follow
- [] Time to build the first 5–10 skills against the standard before the library reaches critical mass
- [] Optional: a local machine for the local-model/proxy infrastructure if cost-controlled local development is a priority
- [] Executive sponsorship — this is a Director-level platform investment, not a single automation

Results

- **40+ skills** shipped and in active use across a 50+-person GTM org.
- Idea-to-scoped-project time cut from **4–8 hours to under 30 minutes** via the project-intake engine.
- Every project in the portfolio traces back to the same build methodology — meaning a new hire (or a hiring committee) can read any one build and understand the standard behind all of them.
- Local AI infrastructure enables zero-API-key experimentation, decoupling exploration cost from production spend.

Challenges & Fixes

- **Every new skill reinventing its own conventions** → a published build spec that every new build inherits from day one.
- **New project ideas taking hours of back-and-forth to scope before any work started** → a self-serve intake engine that turns raw input into a full leadership package in under 30 minutes.
- **Institutional knowledge living in one person's head** → a documented skill library plus a persistent context-engine so work compounds instead of resetting.
- **AI experimentation cost competing with production budget** → a local, self-hosted model stack for exploration, fully decoupled from hosted API spend.



Demonstrates

This is the Director-level story: not automations, but an **operating system for a GTM org's AI adoption** — standards, governance, reusable components, cost control, observability, and a documentation culture, built by one person and used by fifty.

STACK & INTEGRATIONS

Claude Skills/MCP

n8n

Ollama

HubSpot

Clay

Amplemarket

Avoma

Asana

Factors.ai

Microsoft 365

WHAT IT NEEDS

- ✓ AI coding/skill environment access for the org
- ✓ A written intake/PRD convention (e.g. ROSTR or your own equivalent)
- ✓ Time to build the first 5-10 skills against the standard

- ✓ Executive sponsorship — a platform investment, not a single automation

DISCUSS THIS BUILD

Happy to walk through the architecture, the failures along the way, or how this could translate to your team.

Email Patrick



© 2026 Patrick Diamitani. Case studies anonymized for public portfolio use — no employer-identifying or client-identifying detail included.

pdiamitani@atlashxm.com · [linkedin.com/in/diamitani](https://www.linkedin.com/in/diamitani)