

Patrick Diamitani

FORWARD DEPLOYED ENGINEER · GTM AND APPLIED AI SYSTEMS

Chicago, IL (open to relocation) · 319.930.0290 · patrick.diamitani@gmail.com
linkedin.com/in/diamitani · github.com/diamitani · diamitani.com · rostragent.com

SUMMARY

I build enterprise-grade production AI systems that non-technical teams rely on every day. I own AI strategy and automation for the entire go-to-market organization at Atlas HXM, a global Employer of Record platform with about 500 employees across 160+ countries. The production library holds 16 systems, including a two-channel AI inbound system credited with **\$9M+ in attributed pipeline** and **\$16M+ in pipeline enabled** overall.

Alongside that I author and maintain ROSTR, a published architecture for production-grade multi-agent systems, and its open-source implementation. My path here ran through seven years of enterprise sales at Uber, Upwork, and Techstars, which is why I can sit in the discovery call, translate what a customer actually needs, and then go build and deploy it.

Every build ships to the same standard: LLMs handle messy language, deterministic unit-tested code decides anything auditable, and nothing goes live without dry-run tests, a rollback plan, and shared error handling.

TECHNICAL STACK

Languages	Python · JavaScript / TypeScript · Node.js · React & Next.js · SQL and NoSQL · HTML/CSS · Bash
AI & Agents	OpenAI (Agents SDK, Responses API, Assistants, Custom GPTs) · Anthropic Claude (Claude Code, Skills, Agents, Artifacts) · Google Gemini (API, ADK, AI Studio, Vertex) · AWS Bedrock, SageMaker, AgentCore · Azure AI Foundry · Ollama and local inference · DeepSeek · RAG pipelines, evaluation harnesses, prompt architecture
Agent runtimes	ROSTR Framework (author) · Hermes Agent · OpenClaw · Codex · OpenCode · Antigravity (Gemini)
Cloud	AWS: Lambda, S3, IAM, EC2/VPC, API Gateway, RDS, DynamoDB, Redshift, Cognito, Bedrock · Azure: Blob Storage, AI Search, Static Web Apps, Container Apps, App Services, AI Foundry, CI/CD · Supabase · Vercel
Automation	n8n (event-driven multi-agent pipelines) · Make.com · Zapier · HubSpot workflows, webhooks, and custom code · REST APIs and webhooks · Git/GitHub Actions · CLI tooling (AWS, GitHub, Claude, Hermes)
GTM systems	HubSpot · Salesforce · Zoho · Clay · Amplemarket · Apollo · ZoomInfo · Gong · Reply.io · Voiceflow · LinkedIn Sales Navigator

PUBLISHED RESEARCH: THE ROSTR FRAMEWORK

ROSTR: A Unified Architecture for Production-Grade Multi-Agent Systems with Phase-Aware Orchestration and Persistent Knowledge Compounding. Sole author. Published on Zenodo with DOI, April 2026. 22,000 words, 27 references.

The framework defines four pillars for production multi-agent AI. **PAL** is a prompt abstraction layer that compiles natural language into strictly typed agent manifests. **NPAO** is deterministic phase-aware orchestration with human-in-the-loop gates. **RAG DAL** is multi-pass retrieval with source-credibility tiers and contradiction detection. **Rostr**

Hub is persistent cross-session memory.

ROSTR Agent (rostragent.com · github.com/diamitani/rostr-agent) is the MIT-licensed open-source implementation, built on the Hermes runtime with 41 production skills spanning GTM, development, content, data, automation, and operations.

BENCHMARK: 200 TASKS ACROSS 8 DOMAINS, IDENTICAL MODEL, TEMPERATURE, AND TOOLS

+23pp task completion · +28pp first-attempt accuracy · +82pp cross-session knowledge retention · 33% lower token cost per completed task

PROFESSIONAL EXPERIENCE

Atlas HXM | GTM AI & Automation Manager

Chicago, IL · Sep 2025 to Present

Own AI strategy and automation across sales, marketing, and RevOps at a global Employer of Record platform (about 500 employees, 160+ countries). The production library holds 16 systems used daily by 50+ non-technical reps and marketers; my builds have enabled \$16M+ in pipeline to date.

- **Prospect automation engine** — an autonomous five-stage GTM pipeline. (1) Intent signal: ZoomInfo and Factors.ai flags companies researching EOR / global hiring, written to HubSpot daily. (2) Qualification & routing: n8n runs an eight-stage deterministic filter chain (employee count, geography, open deals, lifecycle stage, record type, 30-day recency, competitor suppression, industry) with explicit round-robin territory assignment and 50-state coverage self-check. (3) Enrichment: Clay waterfall people search with verified email, mobile, and secondary phone; waterfall multi-source email verification; company and prospect research via Firecrawl live web search. (4) Personalization: single schema-constrained LLM call generates three email bodies, three subject lines, and a pain-point analysis grounded in research findings, all stored as HubSpot message tokens. (5) Delivery: prospects enrolled in Amplemarket with warmed per-rep dedicated mailboxes; Orum auto-queues call tasks. Architecture evolved across three generations (Nov 2025–Mar 2026): Gen 1 (59 nodes, 3.86% open rate) → Gen 2 (22 nodes, optimized orchestration) → Gen 3 (n8n + Clay, **43–55% open rate, 100% success, 75 companies/week**). See [portfolio](#) for architecture diagrams.
- **Inbound automation engine** — two-channel system analyzing website chatbot and Contact Us form messages. An intent classifier determines whether a submission is a sales inquiry (schedule follow-up), general question (send help resources), press/partnership (route to founder), or spam (silently drop). For sales inquiries, an AI agent drafts a personalized opening message, assigns the prospect to a rep via round-robin, stores the message in HubSpot, and enrolls them in a welcome sequence. Combined with prospect automation, these two systems attributed **\$9M+ in attributed pipeline** and enabled **\$16M+ in total pipeline**.
- Designed and launched an **automated EOR compliance-screening engine** in n8n. An LLM reads messy CRM data and a unit-tested rules engine makes the call. Corrected 100+ misclassified deals, more than doubled detection of ineligible deals, perfect production write record with full audit trail.
- Inherited a failing four-agent post-sales-call pipeline and **rebuilt it from RAG-based design to deterministic architecture**. Zero successful end-to-end runs became fully verified production runs.
- Owned a runaway AI and data-enrichment spend incident, then redesigned around a free eligibility gate and tiered cost profiles. **Spend per record dropped roughly 70%**, enforced by architecture.
- Built and published a **40+ skill AI automation library** plus a self-serve project intake engine. Idea-to-scoped-project time fell from 4 to 8 hours down to under 30 minutes across a 50-person GTM organization.
- Directed a nine-script enrichment pipeline for an outsourced SDR program: confirmed 0% contact overlap against thousands of already-contacted records and 91% phone coverage.

Diamitani Industries, Inc. | Founder, AI Engineer

Remote · Jan 2014 to Present

Active NYC C-Corp holding my consulting practices, product portfolio, and recorded music catalog. About 250 repositories across research, agent frameworks, production apps, and open-source skills.

- Ship and operate 8+ production web applications end to end on Next.js, React, Vercel, Supabase, and LLM APIs from OpenAI, Anthropic, and Google.
- Author and maintainer of the ROSTR Framework and ROSTR Agent, plus **17 published open-source agent skills** covering agent scaffolding, persistent context, n8n execution diagnostics, HubSpot chatbot management, PQL routing, and GTM engineering.
- Consulting entities: **SalesGency** (GTM AI consulting) and **GencyAI** (AI and automation development).
- Recorded music: BMI and SoundExchange member; lo-fi hip-hop and EDM instrumentals released under my own catalog.

GencyAI | AI & Automation Development

Remote · Jan 2024 to Present

An AI development practice for SMB and mid-market clients covering agents, skills, prompt libraries, and production web apps, taken from SOP to shipped system with adoption built in.

- Deployed a production website chatbot for a recruiting agency (Voiceflow, OpenAI) with a custom knowledge base connecting candidates and employers; wrote the PRD and built the supporting Custom GPTs.
- Built a full-stack email outreach platform (v0.dev, Vercel, Supabase, Make.com) that ingests CSVs, processes webhook data, and sends and tracks campaigns for a startup client.
- Rebuilt an education company's site as a custom web app (v0.dev, Vercel, Gemini API) with generative AI course lessons and text-to-speech voice.
- Published **150+ Custom GPTs** in the GPT Store under GPTPAT, spanning programming, productivity, sales, creative, civic, and lifestyle categories.
- Delivered content and commerce automation including a YouTube comment moderation tool, an RSS-driven blog writer matched to the client's voice, and Shopify product-description and transaction-tracking workflows.

Scalytics.io | Founding GTM Engineer, Consulting

Remote · Apr 2025 to Present

- Embedded part-time with the founding team of a private-AI infrastructure startup (PhDs and federated-data experts) to build their go-to-market engine from scratch.
- Developed custom CRM tooling, sales playbooks, and lead-generation systems.
- Integrated Make.com automations with Zoho and HubSpot, and managed pipeline plus LinkedIn marketing.

Dark Roast Media | Sales Enablement Consultant

Remote · Jan 2024 to Dec 2024

- Stood up outbound sales operations end to end with HubSpot, Reply.io, Asana, and LISN. Vetted the tooling, configured automation sequences, and trained the incoming SDR team.
- Identified ICPs and generated 60,000 leads; ran 500 emails a day at a 40% open rate and 12% reply rate.
- Held sales calls with C-suite leaders across venture capital, real estate, and startups.

Rand McNally | GTM Consultant, Sales & Marketing

Remote · May 2023 to Dec 2023

- Created the outbound sales playbook for the Rand Fleet Software launch; configured automation sequences and managed leads for six sales reps, driving a 2,000% increase in engagement.
- Oversaw marketing email campaigns to 200,000+ customers and leads at a 50% open rate.

Uber, Inc. | Enterprise Representative, Uber for Business

Chicago, IL · Aug 2022 to Jan 2023

- Generated \$947K+ in gross bookings managing a book of business across technology, automotive, and retail.
- Led the hiring cohort in quota attainment; top 10% revenue earner in a sales organization of 100+ reps.
- Ran technical discovery and cold outreach with C-suite, director, and VP buyers, averaging 2,000+ emails and 250 dials weekly using ZoomInfo, Gong, and LinkedIn Sales Navigator.

Upwork, Inc. | Senior Enterprise Representative

Chicago, IL · Oct 2019 to Jan 2022

- Top 10% performer, consistently above quota with high call-engagement rates.
- Partnered with Fortune 500 executives to scope strategic use cases for outsourced and distributed talent.

Techstars Chicago | Portfolio Associate, Sales

Chicago, IL · Jul 2018 to Oct 2018

- Supported founders of three accelerator startups on sales infrastructure, playbooks, and go-to-market strategy: 2ndKitchen (acquired), Clyde (acquired), and Rheaply (\$40M raised), over \$100M raised collectively.

DreamItReel | Founding Account Executive

Chicago, IL · Sep 2016 to Jan 2018

- Built the internal sales process from zero and grew revenue from \$200K to \$1.5M+ at 40% year-over-year growth, closing 100+ sales to agencies, Fortune 500 brands, and online publishers while managing a team of five.

Omicon Entertainment | Founding Partner

2010 to 2016

- Produced 1,000+ shows for 80,000 students at \$2M in yearly revenue; trained and managed 30+ interns.

SELECTED TECHNICAL PROJECTS

Founder-built products designed, developed, and deployed end to end, mostly Next.js and React on Vercel with Supabase, plus LLM APIs from OpenAI, Anthropic, and Google.

ROSTR Agent | rostragent.com

Python · Hermes runtime · multi-provider LLM · vector retrieval · MIT license

Self-improving open-source multi-agent framework implementing PAL, NPAO, RAG DAL, and Hub intelligence layers, with 41 production skills. Benchmarked against baseline over 200 tasks in 8 domains. Companion repositories extend the framework to Amazon Bedrock AgentCore, OpenClaw, and a full platform build (backend, core, dashboard, skills).

Artispreneur | artispreneur.com

Next.js 14 · Supabase · Azure AI · Stripe · multi-tenant

SaaS platform giving independent artists the business layer of a record label. Includes an EPK builder, contract generation, catalog management, an academy course platform, and AI assistants. Has supported 100+ artists, and includes a music-industry directory of 78K+ verified contacts.

Replai | replaiall.com

Next.js · LLM rewrite pipeline · per-contact guardrails

AI messaging assistant that rewrites impulsive or emotionally charged texts while preserving the sender's voice, with per-contact rewrite rules.

Lets Vibe AI | letsvibeai.com

Next.js · Vercel · curriculum agent (ROSTR PAL)

Free AI-builder academy with a 6-module curriculum, 3 project labs, a 10-day bootcamp, and an AI curriculum-planning agent that researches learning resources and generates personalized paths.

Lola | trainlola.com

Next.js · DeepSeek · Ollama · ROSTR Framework

AI fitness coach with personalized workout plans, curated classes drawn from social video, progress tracking, and accountability partners.

Huddle CoWork | huddleco.work

Next.js · Supabase · programmatic SEO

Coworking directory for finding, comparing, and booking across 7,000+ locations in 1,100+ US cities.

Civic Pie | [AI civic engagement platform](https://ai-civic-engagement-platform.com)

Next.js · Supabase · RAG over public records

Hyper-local civic platform covering all 50 Chicago wards, with alderman records, meeting data, and a CivicGuide AI assistant. Piloted in the 48th Ward alongside a campaign dashboard, volunteer coordination tool, and constituent feedback-analysis chatbot.

Credit Fixer | fcragent.com

Next.js · LLM guided workflow · FCRA rule set

AI credit repair agent that walks users through identifying and resolving issues on their credit profile.

Pulse AI | Daily AI news aggregator

26 ingestion sources · DeepSeek summarization · RAG chat · Supabase archive

Scheduled ingestion and summarization pipeline with a retrieval-backed chat interface over the archive.

Gency AI | gencyai.com

Next.js · client delivery practice

AI software development practice: agent setup, skills, prompt libraries, and production web apps, fully managed for clients.

OPEN SOURCE & AGENT SKILLS

github.com/diamitani holds about 250 repositories. **17 published agent skills** installable into Claude Code, Hermes, and any ROSTR-compatible runtime:

rostr-agent-builder: scaffolds production multi-agent Python systems from specs

rostr-pal-skill: LLM-agnostic intent compilation into typed manifests

inbound-automation-engine: website chatbot + form intake classifier that routes inquiries and drafts personalized responses to reps via round-robin

n8n-execution-analyst: pulls execution data by workflow ID or link, diagnoses failures, error chains, and fixes; builds observability dashboards

hubspot-ticket-submitter: sends tickets and contact records to HubSpot from Claude or other agents; validates schema and handles async writes

soul-builder: agent personality, values, and behavior constraints

gtm-engineering: GTM system design and build patterns

marketing-chatbot-execution-analyst: pipeline health and conversion analysis

monarch-art-vault: photo folders to print-ready catalogs

context-engine: persistent cross-session memory, indexing, retrieval

prospect-automation-engine: five-stage autonomous lead pipeline (signal → qualify → enrich → personalize → deliver) with live web research, multi-source verification, and per-rep mailbox warmth

n8n-engineer: designs and builds n8n workflows for lead ingestion, enrichment, deterministic multi-stage qualification, and end-to-end automation pipelines

hubspot-import-manager: uploads lead CSVs and configures them to a master template for bulk import to HubSpot; manages custom properties and field mapping

clay-enrichment-orchestrator: manages lead enrichment workflows and waterfall data routing through Clay; sets up multi-table pipelines and webhook triggers

hubspot-chatbot-manager: chatbot config and delivery management

partner-pql-manager: PQL scoring, routing, follow-up in HubSpot

entity-readiness-execution-health: n8n workflow validation and delivery status

curriculum-os: PAL-driven research and personalized curriculum generation

Also maintains **Monarch Workspace** (agentic workspace where agents plan work and pause for approval before anything reaches the outside world), an **AI tool directory** of 900 tools, a public **Custom GPT directory**, AWS troubleshooting guides and templates, and a published **response to the US AI Action Plan request for comments**.

EDUCATION & CERTIFICATIONS

Per Scholas | AWS re/Start Program

Chicago, IL · 2024 to 2025

AWS Certified AI Practitioner (2025) · AWS Certified Cloud Practitioner (2025). Hands-on training deploying servers (VPC, EC2), serverless applications (Lambda, API Gateway), AI/ML services (Bedrock, Polly, Lex), and databases (RDS, DynamoDB, Redshift).

St. John's University, Tobin College of Business | B.S. Economics

Queens, NY

Co-founded the Entrepreneurial Society (200+ members), served in elected Student Government administering a \$1M budget, and participated in the G.L.O.B.E. Microfinance Program.

CIVIC & COMMUNITY

- **48th Ward, Chicago. Precinct Captain.** Ground organizing for a local aldermanic campaign; built the campaign's internal dashboard, precinct coverage and turnout analytics, volunteer coordination tooling, and a constituent feedback-analysis platform.
- **Re:Work Training, Sales Training Volunteer (2022 to Present).** Lead courses and workshops training new BDRs from underrepresented backgrounds on sales fundamentals.