

AI & AUTOMATION · GTM SYSTEMS ENGINEERING

Patrick Diamitani

GTM AI & Automation Manager · AI & Automation lead at a ~500-person global Employer-of-Record (HR-tech) platform.

I build production AI automations that ship — and the systems that let a 50+ person GTM org run on them.

Targeting: Director of AI & Automation · Applied-AI / Forward-Deployed Engineering · Founding GTM-Engineering roles

linkedin.com/in/diamitani

Chicago, IL

Public edition — the employing company and internal names, IDs, and endpoints are anonymized. All results and technical detail are real and verifiable under NDA.

IMPACT AT A GLANCE

8+

production automations live across the GTM stack

~242

pipeline deals auto-screened for EOR compliance, every one auditable

0→stable

post-sales workflow recovered from 21/25 failing to reliable end-to-end

515K

credit spend spike traced to root cause & governed with a reusable template

<30 min

idea → fully-scoped project, down from 4–8 hours

30+

reusable Claude skills & agents shipped for the whole GTM team

TECHNOLOGY & PLATFORMS

ORCHESTRATION

n8n

AI / LLM

Azure OpenAI (GPT-4o / GPT-5-chat)

automation platform that chains apps so steps fire automatically — no human clicking buttons

Webhooks & MCP

event triggers and tool-connectors that let systems call each other in real time

Custom JavaScript engines

deterministic rules written in code, so decisions are repeatable and auditable

GTM DATA & CRM

HubSpot API

the CRM system of record — deals, contacts, properties, workflows

Clay

data-enrichment workspace that builds and enriches prospect lists

Amplemarket & ZoomInfo

contact discovery and firmographic enrichment sources

Avoma

call-recording & transcription intelligence

Factors.ai

first-party website intent & account signals

the reasoning layer that reads messy text and classifies or drafts

Claude + Claude Skills / Agent SDK

the platform used to build the reusable skill & agent library

Retrieval & grounding

feeding models real source data so outputs are specific, not generic

INFRA & DELIVERY

Azure Blob Storage

cloud file store used to serve 250+ country data files deterministically

Microsoft 365 (Outlook, Teams, Excel, OneDrive)

delivery + human-readable mirrors

Onspring

internal governance/compliance data source

PORTFOLIO OF BUILT WORK

FLAGSHIP · LIVE IN PRODUCTION

EOR Compliance Checker (v5)

Auto-screens every sales-pipeline deal against a 30-sector compliance worksheet the moment it enters — and writes an auditable verdict back to the CRM.

n8n · Azure OpenAI GPT-4o · HubSpot API · Avoma API

LIVE · MAJOR RELIABILITY TURNAROUND

Post-Sales Call Automation

After every sales call, reads the transcript and auto-generates a rep-ready follow-up package: consultative email, missed-questions analysis, competitor research, LinkedIn draft, and a logged CRM note.

n8n cloud · Avoma webhook · Azure OpenAI · Azure Blob Storage (250+ country files)

BUILT · REQUESTED BY VP GROWTH

Pre-Sales Call Automation

LIVE · \$250K OPEN DEAL SOURCED

AI Customer Agent — Built End to End

The moment a meeting is booked, sends the rep a concise pre-call prep note so they walk in already knowing where they're most likely to close.

HubSpot workflows · n8n · Factors.ai Journey API · Azure OpenAI

The website AI agent built start to finish — the chat app itself, a 4,100-execution AI qualification pipeline, a three-workflow reporting warehouse, and deal attribution that traced a \$250K open deal back to its originating conversation.

HubSpot chatbot + workflows · n8n · Azure OpenAI · HubSpot CRM API

DEPLOY-READY

Partner PQL Workflow

Classifies inbound partner-form submissions, drafts a partner-specific email, and routes a task to the partnerships owner automatically.

n8n · Azure OpenAI (gpt-5-chat) · HubSpot API

BUILD-READY · APPROVAL-GATED

Entity Readiness Change Report

Every Monday, diffs the company's country hire-readiness against last week and emails Growth only the countries that flipped — as a locked two-column table.

n8n · Onspring · n8n Data Table · Microsoft Excel / OneDrive

ONGOING · OWNS THE PIPELINE

SDR-Agency Prospecting Enablement (AGS)

Runs the full prospecting supply chain for an outsourced SDR agency — clean, enriched, CRM-safe lists activated on a weekly cadence.

Clay · Amplemarket · ZoomInfo · HubSpot API

STRATEGY + BUILD

Outbound Signals & Intent Engine

A defined library of fit + intent signals, a 0–100 account scoring model, and composite “plays” that stack signals into prioritized, repeatable outbound motions.

Clay · Factors.ai · HubSpot · Scoring model (0–100)

SHIPPED · COST + GOVERNANCE WIN

Clay Credit Governance & Cost Control

Traced a runaway enrichment-credit spike to its real sources and shipped a reusable template + guardrails that cut per-row cost and stop the leak repeating.

Clay API · Excel data modeling · PDF one-sheeter for leadership

META-SYSTEM · FRAMEWORK

New Project Automation System (Center of Excellence)

Turns any raw input — notes, files, links, a one-line idea — into a complete, leadership-ready project package: knowledge base, PRD, architecture diagram, reporting framework, and a phased build plan.

Claude / Agent SDK · n8n · Notion · Asana

PLATFORM · BUILT FOR 50+ REPS

GTM Skill & Agent Library (30+)

A library of reusable AI skills and agents that put every automation above into the hands of non-technical reps and marketers through plain-English commands.

Claude Skills / Agent SDK · n8n · Clay · HubSpot

EOR Compliance Checker (v5)

Auto-screens every sales-pipeline deal against a 30-sector compliance worksheet the moment it enters — and writes an auditable verdict back to the CRM.

THE PROBLEM

Reps advanced deals to contracting without knowing if the role or industry was eligible for the service. Compliance only caught blockers at contract review — after the buyer was ready to sign. Lost deals and wasted effort.

WHAT I BUILT

- A 46-node n8n production workflow (real-time webhook + hourly safety-net sweep + on-demand batch re-screen).
- A hybrid engine: the LLM only maps messy CRM data to one of 30 canonical sectors; a deterministic JavaScript matrix makes every tier decision.
- A three-axis compliance framework (viability tier + advisory flags + hard screen-outs) transcribed faithfully from the official worksheet.
- A report-only dry-run harness that tests any future rule change against the entire pipeline with zero writes.

SYSTEMS THINKING

The core design decision — rules live in code, not in the model's head — is what makes it durable. The LLM does only what it's good at (reading unstructured text); every verdict is deterministic, repeatable, and cites its own reasoning. Edge cases were reasoned about explicitly: pass-through detection (a staffing firm placing workers into oil & gas is screened as oil & gas), semantic sector aliasing, and a "needs-title" hold instead of guessing on thin data.

OUTCOME

- Live in production, screening the full sales pipeline (~242 open deals).
- Corrected ~103 deals wrongly stuck in "review required" while tightening genuine ineligibility detection — every flag traceable to a named rule.
- 16 unit tests + full dry-run validation + clean production writes across all deals before go-live.

STACK

n8n

Azure OpenAI GPT-4o

HubSpot API

Avoma API

Custom JS rules engine

NAICS resolver

Outlook

WHAT BROKE & HOW IT WAS FIXED

ISSUE / SYMPTOM	ROOT-CAUSE FIX & PREVENTION
A const-reassignment bug surfaced on the first live production test	Caught immediately by the verification step; fixed and re-tested green. Prevention: the dry-run harness now gates every change.
Five analytics properties silently failed	Folded that content into the audit-note field instead. Prevention:

ISSUE / SYMPTOM	ROOT-CAUSE FIX & PREVENTION
to write (didn't exist in the CRM)	validate property existence before writing.
Old prompt let thin data demote clean sectors	Re-architected so confidence is a report field that never changes the tier.

02

LIVE · MAJOR RELIABILITY TURNAROUND

Post-Sales Call Automation

After every sales call, reads the transcript and auto-generates a rep-ready follow-up package: consultative email, missed-questions analysis, competitor research, LinkedIn draft, and a logged CRM note.

THE PROBLEM

Reps spent time after every call on follow-up research and writing. The automation meant to do it was failing 21 of 25 runs — a dead vector-search index left four AI agents with no country context.

WHAT I BUILT

- An always-on n8n cloud workflow triggered by a call-platform webhook, so it fires regardless of whose laptop is open.
- Four AI agents (follow-up email, missed questions, competitor research, LinkedIn post) grounded in real country data.
- A deterministic replacement for the broken search: extract countries → ISO-3 codes → fetch each file directly from blob storage → inject into every agent.

SYSTEMS THINKING

This is a textbook resilience rebuild. Rather than recreate a fragile index (which needed permissions Patrick didn't have), the whole failure point was designed out: country code → exact file, no fuzzy search, no credentials to rotate, and a graceful fallback so the email still sends if a file is missing. Failure modes were mapped end-to-end — future/scheduled meetings, broken item-lineage, context bloat — and each got an explicit guard.

OUTCOME

- Recovered from 0/25 runs reaching the email step to reliable end-to-end delivery.
- Cut context bloat ~3x (20→6 relevant countries per call) — cheaper and higher signal.
- Restored specific, grounded output over generic filler.

STACK

n8n cloudAvoma webhook

Azure OpenAI

Azure Blob Storage (250+ country files)

HubSpotOutlook

WHAT BROKE & HOW IT WAS FIXED

ISSUE / SYMPTOM	ROOT-CAUSE FIX & PREVENTION
84% of failures: dead Azure AI Search index (DNS not-found)	Removed the dependency entirely; replaced with deterministic blob fetch. Prevention: prefer deterministic lookups over services that can silently die.
Silent dead-ends on scheduled/future meetings	Added a meeting-completion guard + bounded retry budget (10→20 polls) + explicit exit nodes.

ISSUE / SYMPTOM	ROOT-CAUSE FIX & PREVENTION
Broken paired-item lineage dropped the transcript downstream	Carried the base object forward and converted 28 fragile item-references to a stable form.

03 BUILT · REQUESTED BY VP GROWTH

Pre-Sales Call Automation

The moment a meeting is booked, sends the rep a concise pre-call prep note so they walk in already knowing where they're most likely to close.

THE PROBLEM

Reps burned time on manual pre-call research; time-zone-based daily checks risked missing calls booked across regions.

WHAT I BUILT

- HubSpot meeting-booked webhook → n8n → enrichment → one-paragraph prep note emailed to the assigned rep.
- Enrichment pulls first-party intent (Factors Journey API), company data, and the most recent contact message.
- A guardrail against duplicate sends; the summary is pre-built instantly and timed to arrive a set number of hours before the call.

SYSTEMS THINKING

Chose an event trigger (meeting booked) over a scheduled daily sweep specifically to defeat the time-zone failure mode — a systems decision made before the first line was built, and documented with an explicit risk/mitigation in the brief.

OUTCOME

- Scoped, brief-approved, and built against a clear definition-of-ready (clean webhook execution, no hallucinated facts).

STACK

HubSpot workflows n8n
Factors.ai Journey API
Azure OpenAI

WHAT BROKE & HOW IT WAS FIXED

ISSUE / SYMPTOM	ROOT-CAUSE FIX & PREVENTION
Risk: a daily check could miss calls in other time zones	Mitigation designed in up front — trigger on booking, dedupe, and send relative to call time regardless of zone.

04

LIVE · \$250K OPEN DEAL SOURCED

AI Customer Agent — Built End to End

I built the company's website AI agent end to end: the chat application (guidelines, guardrails, knowledge base, routing), the live AI qualification pipeline behind it, the reporting warehouse around it, and the deal-attribution layer that proves what it's worth.

THE PROBLEM

The website chat agent was capturing conversations, but the attribution flag was never stamped — so 364 contacts going back years were invisible to every report. The channel looked dead when it wasn't.

WHAT I BUILT (FOUR LAYERS)

- **The agent:** versioned conversation guidelines & guardrails (v4→v6), knowledge base, qualification and handoff rules — live on the public pricing page. Plus three coordinated HubSpot workflows (chat-end trigger, MQL owner-assignment + notification, non-MQL nurture email).
- **The AI qualification pipeline** (live, 4,100+ logged executions): closed-conversation webhook → contact resolution → hand-written transcript normalizer (visitor-vs-bot actor detection; email/phone/company extraction from free text) → LLM agent with structured-output parsing under governance rules: only visitor messages count as buying intent, confidence < 0.65 forces Non-MQL, blank values never overwrite CRM data. 15 properties stamped per MQL + an AI-drafted follow-up; idempotency flag so nothing double-processes.
- **The reporting warehouse:** hybrid three-workflow architecture — real-time event collector, nightly reconciliation that rebuilds source-of-truth metrics (including a “why didn't this chat create a contact?” gap metric), and an hourly AI-credit watcher with 4-tier alert thresholds and month-end spend forecasting.
- **Deal attribution:** deal → contact → chat-session association walk with three honesty tiers (sourced / influenced / excluded), operationalized as a weekly stamper workflow writing attribution tier + originating conversation ID onto deals. Plus the 364-contact historical backfill (364/364 updated, 0 errors) and five smart lists.

OUTCOME (LIVE CRM, JULY 2026)

- **\$250K open deal sourced by the agent** — chat one day, deal opened the next, traced to the exact originating thread ID (verified in the live CRM, July 2026).
- ~\$520K agent-sourced deal value since launch; ~\$680K closed-won in 2026 on chatbot-engaged contacts (influence tier, labeled as such); ~\$3.6M lifetime deal value across ~390 engaged contacts.
- 364 historical contacts recovered into attribution; idempotent classify → route → notify pipeline live end-to-end.

STACK

HubSpot Customer Agent + workflows

n8n (LangChain agent nodes)

Azure OpenAI

HubSpot Conversations + Associations APIs

Microsoft Graph (Excel warehouse)

Slack notifications

SYSTEMS THINKING

Diagnosed the root attribution failure before building anything (a data audit found the capture-vs-stamp gap), then designed for reliability: UUID session IDs, an idempotency checkpoint, and system-inbox sending so delivery never depends on a person.

WHAT BROKE & HOW IT WAS FIXED

ISSUE / SYMPTOM	ROOT-CAUSE FIX & PREVENTION
Attribution flag captured chat data but never stamped the source	Built a stamp workflow for go-forward + API backfill for history. Prevention: audit the whole capture→report chain, not just capture.

05

DEPLOY-READY

Partner PQL Workflow

Classifies inbound partner-form submissions, drafts a partner-specific email, and routes a task to the partnerships owner automatically.

THE PROBLEM

Partner inquiries needed manual triage, classification, and routing before anyone could act.

WHAT I BUILT

- An n8n workflow: form webhook → AI classify (partner vs. not) → AI-drafted partner email → lifecycle set → CRM task routed to the partnerships owner.
- Hardened for production: fixed a v3-API association bug (id vs. vid) that would have created tasks with no contact link, and wired a shared error-handler workflow.

SYSTEMS THINKING

Explicitly matched the error-handling pattern of the existing inbound workflow so failures surface in one place — thinking about the operational estate, not just the single workflow.

OUTCOME

- Deploy-ready workflow; association bug fixed and error handling wired. Gated only on upstream CRM-property + form dependencies.

STACK

- n8n
- Azure OpenAI (gpt-5-chat)
- HubSpot API

WHAT BROKE & HOW IT WAS FIXED

ISSUE / SYMPTOM	ROOT-CAUSE FIX & PREVENTION
HubSpot v2.2 nodes use the v3 API, which returns id not vid	Read id first with a vid fallback, so tasks always associate to the contact. Prevention: verify API-version field shapes on import.

06

BUILD-READY · APPROVAL-GATED

Entity Readiness Change Report

Every Monday, diffs the company's country hire-readiness against last week and emails Growth only the countries that flipped — as a locked two-column table.

THE PROBLEM

Nobody caught weekly changes in which countries the company could legally operate in. Marketing pays for country-specific search ads and didn't want to keep bidding on a country it could no longer serve — or miss one that just opened.

WHAT I BUILT

- An importable n8n workflow: weekly trigger → pull readiness → normalize to a boolean → diff vs. saved baseline → report only the flips → deliver → then write the new baseline.
- A hard no-cloud constraint honored: state lives in an n8n Data Table (source of truth) mirrored to an Excel file for humans — no Azure.

SYSTEMS THINKING

Two constraints drove the architecture and were respected exactly: baseline writes happen only after a successful, non-empty pull (a failed run never corrupts last week's state), and the output stays locked to two columns per the stakeholder's requirement. First-run seeding, empty-pull, and failed-pull edge cases are each handled.

OUTCOME

- Build-ready spec + importable workflow delivered; routed for leadership sign-off. Phase-2 external feed built but held pending a format spec.

STACK

n8n

Onspring

n8n Data Table

Microsoft Excel / OneDrive

Outlook / Teams

WHAT BROKE & HOW IT WAS FIXED

ISSUE / SYMPTOM	ROOT-CAUSE FIX & PREVENTION
Risk: a failed or empty data pull could overwrite good state	Baseline writes gated behind a successful non-empty pull; normalize step throws on empty. Prevention: never let a failed run mutate the source of truth.

07

ONGOING · OWNS THE PIPELINE

SDR-Agency Prospecting Enablement (AGS)

Runs the full prospecting supply chain for an outsourced SDR agency — clean, enriched, CRM-safe lists activated on a weekly cadence.

THE PROBLEM

A paid sales agency needed a steady stream of clean, targeted, CRM-safe lists — without contacting existing customers, open deals, closed-lost, or recently-touched contacts.

WHAT I BUILT

- A repeatable intake → normalize → dedupe → enrich → suppress → activate pipeline across Clay, Amplemarket, ZoomInfo and HubSpot.
- Live HubSpot suppression screening for customer status, open deals, and closed-lost on every company before send.
- A competitive-takeaway play: 188 verified HR contacts at companies currently on a named competitor EOR, ~100 net-new to CRM.

SYSTEMS THINKING

Treats prospecting as a governed data pipeline with source-of-truth boundaries (CRM = truth, enrichment = workspace, orchestration = glue) and a compliance-first suppression layer — the difference between a list and a liability.

OUTCOME

- Thousands of contacts activated; weekly batches tagged by segment for agency-level reporting.
- Every list screened live against the CRM before send (open-deal, closed-lost, active-sequence, recent-touch suppression).

STACK

Clay

Amplemarket

ZoomInfo

HubSpot API

WHAT BROKE & HOW IT WAS FIXED

ISSUE / SYMPTOM	ROOT-CAUSE FIX & PREVENTION
Parallel one-off workbooks caused drift and duplicated effort	Consolidated into one master workbook + segment tabs as the operating source of truth.

08

STRATEGY + BUILD

Outbound Signals & Intent Engine

A defined library of fit + intent signals, a 0–100 account scoring model, and composite “plays” that stack signals into prioritized, repeatable outbound motions.

THE PROBLEM

Broad, untargeted outbound was expensive and easy to ignore, with no defensible reason for any given touch.

WHAT I BUILT

- Eight signal categories (expansion, funding, hiring, leadership, competitor/tech, compliance, first-party engagement, suppression), each weighted and mapped to a data source, persona, and query logic.
- A flagship “Displacement” play targeting Director+ People leaders who hire internationally and already use a competitor EOR.
- A companion Master Signal & Intent Library workbook feeding Clay build sessions.

SYSTEMS THINKING

Turns intuition into a system: every touch has a named, weighted, time-bound reason behind it, which is what lifts reply rates and keeps enrichment credit and rep time focused on accounts with a real reason to buy now.

OUTCOME

- Framework + scoring model + flagship play documented for GTM leadership and operationalized in the enrichment stack.

STACK

Clay

Factors.ai

HubSpot

Scoring model (0–100)

09

SHIPPED · COST + GOVERNANCE WIN

Clay Credit Governance & Cost Control

Traced a runaway enrichment-credit spike to its real sources and shipped a reusable template + guardrails that cut per-row cost and stop the leak repeating.

THE PROBLEM

Enrichment credit usage spiked from <15K/month to ~515K in a single month. ~21% of a month's spend ran through unattributed duplicate tables, phone numbers were bought 2–3x per contact, and searches ran with no fit gate.

WHAT I BUILT

- A full-year forensic audit (a seven-sheet master dataset) that re-attributed the unnamed “Copy of” spend back to its real source tables.
- A Standard Enrichment Table Template: a free ICP-fit gate in front of every paid column, a 3-column deduped phone model capped at 3 attempts, and AI/HTTP columns gated to top-tier accounts.
- Three build profiles (Lean ~15 / Balanced ~21 / Full ~50 credits per row) replacing an undisciplined ~50-credit average.

SYSTEMS THINKING

A credit-planning architecture, not a one-time cleanup: gate → cap → clone-don't-copy → dedupe → consolidate → review weekly. The weekly review replaces after-the-fact audits so the next spike is caught in days, not months.

OUTCOME

- Root cause identified and a leadership one-sheeter + master dataset delivered, replacing ad-hoc exports.
- Per-row cost brought under explicit control; a standing weekly credit review installed as the guardrail.

STACK

Clay API

Excel data modeling

PDF one-sheeter for leadership

WHAT BROKE & HOW IT WAS FIXED

ISSUE / SYMPTOM	ROOT-CAUSE FIX & PREVENTION
A single day burned 201,924 credits from a phone-waterfall re-buying the same numbers	Split phone into three deduped columns capped at 3 provider attempts. Prevention: dedupe + cap before any waterfall runs.

10

META-SYSTEM · FRAMEWORK

New Project Automation System (Center of Excellence)

Turns any raw input — notes, files, links, a one-line idea — into a complete, leadership-ready project package: knowledge base, PRD, architecture diagram, reporting framework, and a phased build plan.

THE PROBLEM

Projects were scattered across a task tool (inconsistent), chat threads (buried), and nowhere (lost). Leadership had no clean answer to “what are we building, why, and when.” Every project was an island only one person understood.

WHAT I BUILT

- An intake engine built on a documented framework (PAL for depth, NPAO for work classification, 4Ds for phasing, RAG/DAL for provenance) that outputs a standardized package every time.
- A searchable project index/database so any initiative is one link away for leadership and onboarding.
- The engine that spawned the reusable skills below — projects come out clean, complete, and consistent.

SYSTEMS THINKING

This is the systems-thinking artifact of the whole portfolio: a repeatable operating model that standardizes how an entire team scopes, documents, and ships work — designed so quality is structural, not dependent on any one person.

OUTCOME

- Idea → fully-scoped project cut from 4–8 hours to under 30 minutes.
- A permanent, searchable Center of Excellence with consistent docs across every initiative.

STACK

Claude / Agent SDK

n8n

Notion

Asana

RAG web search

GTM Skill & Agent Library (30+)

A library of reusable AI skills and agents that put every automation above into the hands of non-technical reps and marketers through plain-English commands.

THE PROBLEM

Powerful automations are useless if only the builder can run them. A 50+ person GTM org needed self-serve, rep-friendly tools — not one-off scripts.

WHAT I BUILT

- Prospecting & enrichment: Clay prospecting engine, credit estimator, CSV enrichers, HubSpot import formatter, list routers.
- Workflow & ops: n8n workflow architect, execution analyst, GTM architect, project intake & handoff, close-out reporting.
- Revenue & enablement: HubSpot data/revenue analysts, RFP knowledge-base builder + answer retriever, Factors.ai signal operator, case-study builder.
- Content at scale: a full branded content suite (email, social, ads, blog, decks, one-pagers) plus an agent-factory that builds new skills from a description.

SYSTEMS THINKING

A platform play: reusable skills with clear triggers, composability (skills call specialist skills), and guardrails baked in (secret-handling, CRM-safety, approval gates) — so scaling to 50+ users doesn't scale the risk.

OUTCOME

- 30+ skills/agents shipped and organized into installable suites for the GTM team.
- Repeatable, governed self-service replacing builder-dependent one-offs.

STACK

Claude Skills / Agent SDK n8n

Clay HubSpot Factors.ai

Avoma

Patrick Diamitani

GTM AI & Automation Manager · Chicago, IL

Contact via LinkedIn: [linkedin.com/in/diamitani](https://www.linkedin.com/in/diamitani)

Every metric and technical detail on this page is drawn from real, shipped project records. Employer and internal identifiers anonymized for public sharing. No secrets, keys, tokens, or endpoints are reproduced anywhere in this document.