

\$9M in Attributed Pipeline Since Inception — An AI Inbound Lead Automation, Built and Proven End-to-End

Case Study / Project Overview — FULL VERSION (internal) · Patrick Diamitani, GTM AI & Automation Manager, Atlas HXM · July 2026

1. At a glance

- **The challenge:** Atlas HXM's inbound "Contact Us" leads sat untriaged — no instant qualification, no routing, and no way to prove what marketing automation was worth in dollars.
- **What I built:** An AI-powered n8n workflow that reads every inbound form submission, classifies intent with Azure OpenAI, sets MQL lifecycle stages in HubSpot, routes partnerships, and replies to sales-ready leads automatically — plus a measurement pipeline that ties every run to CRM revenue.
- **The result:** Across its complete production history since inception: **101 MQLs auto-created (72% of runs), 29 sales deals, \$6.38M in attributed pipeline, \$130K closed-won, and \$9.06M total estimated impact** — every dollar traceable to a specific workflow execution ID.
- **Stack at a glance:** n8n · Azure OpenAI (GPT agents) · Azure AI Search (vector store) · HubSpot CRM + API · Microsoft Outlook · Claude (agentic analysis) · HubSpot MCP · Python · Excel

2. The challenge

Atlas HXM is a global Employer of Record (EOR) platform — it legally employs workers in 160+ countries on behalf of client companies. Its "Contact Us" form is a front door for everything: real buyers asking to speak with sales, job seekers, employment-verification requests, press, and partnership pitches, all landing in the same queue.

That mix created two expensive problems. First, sales-ready leads waited in line behind noise — a buyer asking for "EOR services in South Africa" got the same treatment as a résumé drop. Second, and more damaging for the automation program: nobody could answer the question *"what has this actually made us?"* Marketing automation without revenue attribution is a cost center in every budget conversation.

The project had two halves: build the automation, then prove its economic value with numbers a CFO would accept.

3. The approach

The design principle was **classify once, act everywhere**. Instead of routing raw form text to humans, an AI agent reads each submission and makes one decision — what is this person actually asking for? — and every downstream action (CRM update, lifecycle stage, partner routing, auto-reply) flows from that single classification.

For measurement, the standard shortcut would be a HubSpot attribution report. I rejected that in favor of **execution-level attribution**: pull every individual run of the workflow from the n8n API, extract the exact contact that came through and what the workflow did to them, then match those contacts against CRM deals with a strict rule — *a deal counts only if it was created on or after the day that contact came through the workflow*. This deliberately excludes pre-existing deals on returning contacts (including a \$337K closed-won deal that predated the workflow by one day), making the final number conservative and defensible.

4. What I built (how it worked)

Navigate. Mapped the inbound funnel end to end: HubSpot form submission webhooks → what categories of inquiry actually arrive (seven, it turned out) → which HubSpot lifecycle mechanics mark a contact as an MQL → where partnership inquiries need to land.


Prioritize. The "Speak with Sales" path first — it's where revenue lives. Career inquiries, verification requests, and press get classified and tagged but consume no sales time.

Allocate. The 18-node n8n workflow divides the work:

- A **webhook node** receives every HubSpot form submission in real time.
- An **Azure OpenAI "Analyze Category" agent** classifies the message into one of seven intents (Speak with Sales, Career Enquiry, Employment Verification, Partnership Enquiry, General Enquiry, Press Enquiry, Customer Support).
- A **code node** applies business logic: should this contact's lifecycle stage move to Marketing Qualified Lead? Is this a partnership that routes elsewhere?
- **HubSpot nodes** write the classification, update the lifecycle stage, and log notes — the CRM is always the source of truth.
- A second **AI agent, grounded by an Azure AI Search vector store** (retrieval-augmented generation over Atlas knowledge), drafts a contextual reply to sales-ready leads — referencing the country and service the prospect asked about.
- A **partner path** creates a HubSpot note and sends a routed email via Microsoft Outlook when the AI detects a partnership inquiry.

Orchestrate — the measurement pipeline. Once the workflow was in production, I built the proof layer using Claude as an agentic analyst:

1. A resumable Python client pulled **the workflow's complete execution history (166 runs)** from the n8n API (read-only, cursor-paginated, time-boxed batches to survive sandbox limits).
2. A parser walked each execution's node-level output data and extracted the contact (name, email, company, message), the AI's classification, whether the MQL node fired, and any errors.
3. The **HubSpot MCP** (an AI-to-CRM connector) matched all 121 unique emails against live CRM records — 116 found — then pulled every deal on those contacts with amounts, stages, close dates, and WSE counts (workers to be employed through Atlas).
4. Attribution logic separated the 29 workflow-attributed deals from 5 pre-existing ones, and a formula-driven Excel workbook made every number auditable down to the execution ID.



5. Tools & integrations

Tool / Tech	What it is (plain English)	How it was used here	Connection / auth
n8n	Automation platform that chains apps so steps fire without humans	Runs the 18-node inbound workflow; its API supplied the execution history	n8n API key (kept in environment variables, never in code)
Azure OpenAI	Microsoft-hosted GPT models	Two AI agents: intent classification + drafting grounded sales replies	Azure credentials managed in n8n
Azure AI Search	A vector database — lets AI search company knowledge by meaning	Grounds the reply agent in real Atlas content (RAG)	Azure credentials managed in n8n
HubSpot CRM + API	The CRM and system of record	Lifecycle updates, notes, and the deal/revenue data for attribution	Scoped app token via MCP (value never exposed)
HubSpot MCP	Model Context Protocol — a secure connector that lets an AI operate inside live tools	Let Claude query contacts, deals, and properties directly	Scoped OAuth connection
Microsoft Outlook	Email	Sends routed partnership emails	OAuth via n8n

Tool / Tech	What it is (plain English)	How it was used here	Connection / auth
Claude (Anthropic)	Agentic AI model	Ran the entire measurement pipeline: API ingestion, parsing, CRM matching, workbook build	—
Python + openpyxl	Scripting + Excel generation	Resumable API client, execution parser, formula-driven workbook	—

6. Systems thinking

- **Attribution integrity over headline size.** The rule "deal created on/after submission" cost the headline number over \$1.1M in pre-existing deals — including a \$337K win that closed one day before the workflow's first run. The smaller number survives due diligence; an inflated one doesn't.
- **Failure-aware ingestion.** The n8n pull was built as a resumable, cursor-checkpointed client because the analysis sandbox kills long-running processes. Every batch persisted state; any interruption resumed cleanly instead of restarting.
- **Test-data hygiene.** 22 of 166 runs were internal test submissions (flagged by pattern). They're visible and labeled in the model rather than silently deleted — transparency beats tidy data.
- **Honest gaps, priced anyway.** 9 attributed deals had no amount in HubSpot yet. Rather than ignore them (understates) or guess (fabricates), the model derives a **median value-per-WSE (\$62,453) from the 18 attributed deals that do have amounts**, applies it to the 43 estimated WSEs on the no-amount deals, and exposes the rate as a single editable assumption cell.
- **Everything cites a source.** Every contact row carries its n8n execution ID; every deal its HubSpot record ID. Any number in the summary can be traced to a raw API record.

7. Results & impact

Complete production history since inception (first live run Jun 18, 2026 → Jul 18, 2026):

Metric	Value
Form submissions processed	166 (140 successful runs)
MQLs created automatically	101 — 72% of successful runs
Sales-ready leads answered by AI in real time	112
Partnership inquiries auto-routed	12
Deals attributed to the workflow	29
Attributed pipeline (deals with amounts)	\$6,377,441
Closed-won revenue already banked	\$130,324 (Hill Brothers – Mexico; Great Plains – France)
Open attributed pipeline	\$5,815,117
Estimated value of 9 no-amount deals (43 WSEs × \$62,453)	\$2,685,479
Total estimated attributed impact	\$9,062,920
Workers (WSEs) on attributed deals	122+

Largest attributed deals: Carhartt (\$2.0M, 47 WSEs), Latitude AI (\$1.7M), Arjun Infrastructure (\$664K), Primero AI (\$432K).

8. What this demonstrates

- **Full-lifecycle automation ownership** — designed, shipped, and then *financially proved* an AI system, closing the loop most automation projects leave open.
- **Multi-system integration** — n8n, two Azure AI services, HubSpot API/MCP, and Outlook orchestrated into one flow; then an agentic AI pipeline layered on top for measurement.
- **Revenue-grade data rigor** — conservative attribution rules, test-data flagging, provenance to source-record IDs, and transparent assumptions a finance team can audit and adjust.
- **AI applied where it pays** — LLM classification and RAG-grounded replies aimed at the highest-value moment in the funnel: the minutes after a buyer raises their hand.

Sources: n8n API (workflow 7Exlq2OSVUSh0Tni — all 166 executions since the workflow's first production run on 2026-06-18), HubSpot CRM portal 20072142, pulled 2026-07-18. Full data model: MQL_Inbound_Automation_Economic_Impact.xlsx.

Generated 2026-07-18 · Data: n8n execution history + CRM records · All figures verified against source records.